



Computer Vision

Segmentation – Classic Methods

Technische Hochschule Rosenheim
Winter 2024/25
Prof. Dr. Jochen Schmidt



images: Li, Karpathy, Johnson, CS231n,
lecture 8, Winter 15/16, Stanford

- Contours
 - Canny edge detection
 - Split & Merge
 - Hough transformation
 - Straight lines
 - Circles
- Regions
 - Split & merge for regions

Result of the segmentation: *Segmentation objects* with

- Attributes
- Relationships with each other
- Examples of segmentation objects:
 - Lines
 - Regions
 - Vertices
 - in 3D: Surfaces and volumes
 - In video sequences: moving objects
- Examples of attributes:
 - Position in the image
(qualitative – in front of, behind, ...
quantitative – coordinates)
 - Position in world coordinates
 - Size
 - Quality
 - Gray value, color or texture
 - Motion (2D or 3D;
qualitative – from top to bottom;
quantitative – with x pixel/s, y m/s,...)
 - Shape (qualitative – round, rectangular;
quantitative – gradient,
surface normal)

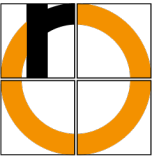
- Segmentation is based on **changes in the image**
(there is nothing to segment in an image with a constant gray value)
- But: Changes are only *indications* of segment boundaries, i.e. they are generally
 - not necessary (1) and
 - not sufficient (2) for segment boundaries
- Examples:
 - (1) 2 objects of the same color located directly adjacent to each other
 - (2) Gray value in the image can change due to shadows/lighting without a segment boundary

In general:

- Decide for one or more parameters (e.g. gray value)
- Track changes of these parameters in the image

Two basic approaches to segmentation


- Find **changes**
 - segment border if change in parameters is sufficiently large
- Track **homogeneous areas**
 - Assumption: homogeneous area, if parameters change only slightly



Objects against a homogeneous background:
Thresholding method Chapter “Preprocessing”

- Calculate threshold value from histogram
- Binarization of the image
- Operations on binarized image (e.g. opening, closing)
- Result: Separation into foreground (object) and background pixels





Contour Lines

Contour line: not necessarily straight, can be curved



Typical procedure:

- **Filter** (gray) image with edge detector
 - e.g. Sobel, Laplace, ... see previous chapters
- **Threshold** value on edge image:
only consider edges above a certain strength → binary edge image
- Optional: Determine the points on the contour line
- **Approximation** of the contour
 - through straight lines
 - piecewise linear
 - circles/ellipses
 - other parametric curves


- Criteria for optimal edge detector
 - Good detection
 - Filter responds to edges, not noise
 - No existing edges are overlooked
 - No non-existent edges are detected
 - Good localization
 - Edges should be detected exactly where they appear in the image
 - Single answer
 - the detector should only provide one response per edge (no double edges)
- Procedure
 - Smooth the image with a 5x5 Gaussian filter (removes edges that are too small)
 - Calculate the gradient (strength and direction)
 - Non-maxima suppression (localize the best edge position)
 - Hysteresis threshold method for edge tracing

Canny – Smoothing / Gradient


$\sigma = 1.0$



$\sigma = 2.0$



$\sigma = 4.0$






Non-maxima suppression

Where exactly is the edge?


- check edge strength along the gradient
- only use the position with the highest value

Technically: Interpolation of the gradient strength required




- Strength of real pixel
- Strength of interpolated pixel

Suppress p if $p < p_{12}$ or $p < p_{34}$
(set p to zero)




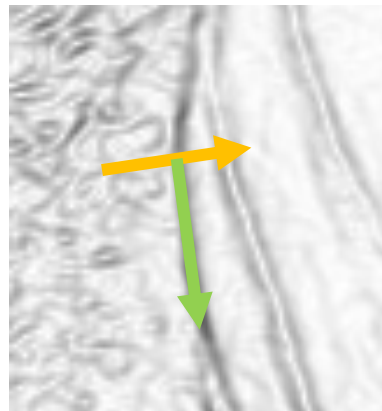
In practice: Instead of interpolation → use discrete directions in 8-neighborhood

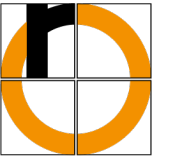


Suppress if the strength is less than one of the two marked neighbors

Hysteresis threshold method

- go through the complete thinned edge strength image
- start a new contour line, if the strength is greater than a threshold value h
- Determine the next edge pixel
 - tangential to the edge, i.e. go perpendicular to the gradient
 - Continue to follow the contour, if the strength is greater than a threshold value l .
 - It applies: $l < h$
- save the contour line created in this way (e.g. as chain code or binary image)





- Either
 - 4-Neighborhood or
 - 8-Neighborhood
- Representation of a contour line as a chain code

	3	
2	+	0
	1	


5	6	7
4	+	0
3	2	1

(x, y)

Coordinates of start


r_1, r_2, \dots, r_n

Directions of the following edge pixels



$(4, 2), 0, 1, 2, 3, 1, 0, 4, 3, 3, 1, 1, \dots$

Canny – Examples



$l = 0,1 \quad h = 0,3$

$l = 0,2 \quad h = 0,3$

$l = 0,1 \quad h = 0,4$


$l = 0,2 \quad h = 0,4$

$l = 0,2 \quad h = 0,6$

$\sigma = 1,0$



$\sigma = 4,0$



Objective: piecewise linear approximation of the contour line

Given:

- Set of (**ordered**) contour points P
- Residual error measure for determining the approximation quality
- Threshold value θ for maximum permissible error

Steps:

1. Initialization
2. Split
3. Merge
4. Shift

- Start with an **initial approximation** – piecewise linear
- Straight line segments are defined by their start and end points, which must lie on the contour
- Initialization in principle arbitrary, e.g.
 - two adjacent pixels form a line segment
 - only the start and end points of the contour form a line segment
 - Every x^{th} contour point is the start/end of a straight line segment

- For each line segment, check whether the error between the approximation and the contour $> \theta$
 - no: segment remains as it is
 - yes: **segment is split** into two parts
- Iterate until error for all pieces $< \theta$
- Split e.g.
 - in the center
 - at the point with the largest error
- This process *increases* the number of straight line segments

- For each pair of adjacent line segments, check whether the merged segment still has an approximation error $< \theta$
 - no: segments remain as they are
 - yes: **segments are merged**
- Iteration, as long as merging is possible
- This process *reduces* the number of straight line segments

- Consider **neighboring pairs** of straight line segments s_1 and s_2 .
Let the corresponding approximation error be ε_1 and ε_2 .
- **Move the end point** of s_1 (= start of s_2) one pixel to the left or right and calculate the errors ε_1' and ε_2' or ε_1'' and ε_2'' .
- Select the contour pixel that leads to the smallest overall error as the new end/start point.
- Iterate until no more end/start points are changed
- This process reduces the error with a *fixed* number of straight line segments

- Since vertical line segments can occur: Use of the Hesse normal form
- Given: Start/end point of the line $(x_S, y_S), (x_E, y_E)$
- Line equation: $ax + by = c$
with

$$a = (y_S - y_E), \quad b = (x_E - x_S), \quad c = y_S(x_E - x_S) - x_S(y_E - y_S)$$

- Vertical distance d_i of a point (x_i, y_i) from the line:

$$d_i = \frac{ax_i + by_i - c}{\sqrt{a^2 + b^2}}$$

Residual error measures for a straight line segment (N : number of points)

- Average distance (L_1 -Norm)

$$\varepsilon_1 = \frac{1}{N} \sum_i |d_i|$$

- Mean square distance (L_2 -Norm)

$$\varepsilon_2 = \frac{1}{N} \sum_i d_i^2$$

- Maximum distance (L_∞ -Norm)

$$\varepsilon_\infty = \max_i |d_i|$$

- approximates contours using parametric functions (e.g. straight lines, circles)
- Suitable for functions with 2 - 3 unknown parameters
- No ordered point set required, no contour tracing necessary




- Line equation in polar form

$$r = x \sin \varphi + y \cos \varphi$$

Two points of view:


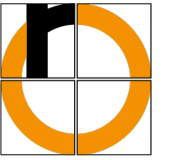
- “Normal” viewpoint:
 - y is the dependent variable, x the independent variable
 - returns a straight line in (x, y) -space for fixed parameters φ_i, r_i
- Dual viewpoint:
 - r is the dependent variable, φ the independent variable
 - returns a harmonic function in (φ, r) -space for fixed parameters x_i, y_i



Idea:

- Each contour point (x_i, y_i) is assigned a harmonic function in (φ, r) -space
- If 3 or more contour points lie on a straight line with the parameters φ_i, r_i , then all their harmonic functions intersect at this one point
- Notes:
 - Of course, this also applies to 2 points – so why look only at 3 or more?
 - In practice, the edge points will never lie exactly on a straight line → clusters in Hough space
 - Low-pass filter
 - Search maximum in a neighborhood




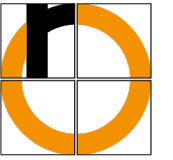
Hough Transform – Straight Lines



Daf-de, [Hough-example-result-en](#), [CC BY-SA 3.0](#)

- Quantization of the parameter space
 $\varphi_i, i = 1, \dots, p$ $r_j, j = 1, \dots, q$
- Initialization Hough accumulator $H(i, j)$ with zeros
- For all edge pixels $(x_k, y_k), k = 1, \dots, N$
 - For all discrete values of $\varphi_i, i = 1, \dots, p$
 - $r_j = x_k \cos \varphi_i + y_k \sin \varphi_i$
 - $H(i, j) = H(i, j) + 1$
- The element $H(i, j)$ of the accumulator contains the number of edge points on the straight line with the parameters (φ_i, r_j)

Hough Transform – Straight Lines




- the Hough transform calculates the parameters of a line equation: the start and end points of the line segment are unknown
- Searching in parameter space is very time-consuming
 - therefore, this should have a small dimension
 - better efficiency e.g.
 - through resolution pyramid
 - use knowledge about the edge direction (restricts the search space)
- Parameter space must be finite
 - An array is created for the accumulator

- Circle detection in principle works in the same way as for straight lines
- Parametric function: Circle equation

$$(x - x_c)^2 + (y - y_c)^2 = r^2$$

- Problem:
 - Hough accumulator becomes three-dimensional
 - Computational effort very high
- Idea: Use the edge gradient

- Consider straight lines on circle contour perpendicular to gradient direction
- all straight lines intersect at the center of the circle
(for discrete noisy images: at least approximately)



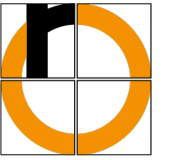
- Quantization of the parameter space (x_{ci}, y_{cj}, r_k)
 $i = 1, \dots, M_x$ $j = 1, \dots, M_y$ $k = 1, \dots, p$
- Initialization Hough accumulator $H(i, j, k)$ with zeros
- For all edge pixels (x_l, y_l) , $l = 1, \dots, N$
 - Calculate normal vector \mathbf{n}_l from gradient direction
 - For all discrete values of r_k $k = 1, \dots, p$

$$\begin{pmatrix} x_{ci} \\ y_{cj} \end{pmatrix} = \begin{pmatrix} x_l \\ y_l \end{pmatrix} \pm r_k \mathbf{n}_l$$

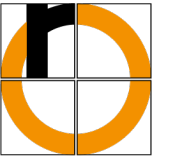
- $H(i, j, k) = H(i, j, k) + 1$
- Search for accumulations in the accumulator $H(i, j, k)$



- the quantization for the radius is typically chosen very coarsely, e.g., 4 to 10 different radii
- Instead of two points (x_{ci}, y_{cj}) for each radius, it can also be useful to enter two whole line segments (as not all lines intersect exactly at the center point)
- As with straight lines, a resolution pyramid is also helpful here

Hough Transform – Circles





Hough Transform – Circles





- (1) closed contour is broken up
- (2) a line is not found
- (3) another line is found in the vicinity of a line
- (4) a non-existent line is found
- (5) line direction and/or length are incorrect
- (6) corners are not detected correctly
- (7) one line is split into several
- (8) an arc with a large radius is detected as a straight line (or vice versa)





Regions

- a region is a homogeneous image area
- a region is always surrounded by a closed contour
segmentation of contours does not always provide a closed contour
- for segmentation based on gray values (simplified):
 - Detection of contour based on sufficiently strong gray value change
 - Region has almost constant gray value
- Segmentation can be done, e.g., based on
 - Gray value
 - Color
 - Texture

- image is broken down into contiguous subsets

$$\mathbf{f} \rightarrow \{\mathbf{f}_i \mid i = 1, \dots, n\}$$

- each pixel belongs to exactly one region

$$\bigcup_{i=1}^n \mathbf{f}_i = \mathbf{f}, \quad \mathbf{f}_i \cap \mathbf{f}_j = \emptyset \quad \forall i \neq j$$

- there is an evaluation criterion for homogeneity

$$H(\mathbf{f}_i) = \begin{cases} 1 & \text{if region satisfies the criterion} \\ 0 & \text{otherwise} \end{cases}$$

- all regions are homogeneous in the sense defined above

$$H(\mathbf{f}_i) = 1 \text{ for } i = 1, \dots, n$$

- all regions are as large as possible

$$H(\mathbf{f}_i \cup \mathbf{f}_j) = 0 \text{ for } i \neq j \text{ and adjacent } \mathbf{f}_i, \mathbf{f}_j$$

- Based on the gray values in the region

$$H(f_i) = \begin{cases} 1 & \text{if } |f_{\max} - f_{\min}| < \theta \\ 0 & \text{otherwise} \end{cases}$$

f_{\max} and f_{\min} are the largest/smallest gray value in the region under consideration (pre-filtering required)

- The threshold value $\theta(a)$ increases linearly with the area a of the region



$$\theta(a) = \theta_S - \frac{a}{a_{\max}} (\theta_S - \theta_E)$$

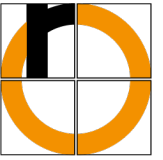
- similar criterion for color (multiple channels) possible

- Split
- Merge
- Split & Merge


- start with a single region (i.e. the entire image)
- normally the image is not homogeneous → further decomposition
 - randomly into two regions
 - regularly, into four equally sized regions (quadtree)
 - ...
- repeat until all regions are homogeneous
- obviously: at some point you always get homogeneous regions, no matter how you decompose them
- the objective is to obtain the largest possible regions as quickly as possible

Split – Quadtree





- Start with a large number of small regions (e.g.: each pixel is a region)
- merge neighboring regions as long as the resulting region is still homogeneous
- This is also the principle of “Selective Search” (segmentation algorithm as the basis of R-CNN)
 - uses a fast graph-based segmentation method as starting regions
 - works (if desired) in several color spaces in parallel
 - 4 homogeneity criteria, which are combined
 - Color (based on histograms of the individual channels)
 - Texture
 - Size of a region (penalizes regions that are too small)
 - Avoiding holes: how well does one region fit into another?



Uijlings, J.R.R., van de Sande, K.E.A., Gevers, T., Smeulders, A.W.M. Selective Search for Object Recognition. *Int. J. of Computer Vision* **104**, 154-171 (2013).
<http://huppelen.nl/publications/selectiveSearchDraft.pdf>

- Principle like split & merge for contours
- Start with an initial decomposition of the image that corresponds to a complete quadtree up to a defined level
- Check all quadrants of the image
 - **Merge** all four regions if H remains fulfilled
 - until merging is no longer possible
- Check all regions
 - **Split** a region into four parts if H is not fulfilled
 - until splitting is no longer possible
- **Group** neighboring regions
 - as long as H remains fulfilled (no longer in the quadtree)
- **Eliminate** regions that are too small
 - if the area of a region is too small: merge it with the neighboring region that differs least from the region that is too small

Contours

- Edge images, e.g., Canny
- Split & merge algorithm
 - Piecewise linear approximation
 - Ordered point set required
- Hough transform
 - Straight lines
 - Circles
 - No ordered point set required

Regions

- Region always has a closed contour line
- Split and/or merge for regions
 - e.g., based on Quadtree
 - Requires
 - Initial decomposition
 - Homogeneity criterion
 - Threshold value(s) for splitting/merging